

```
student@student-HP-Pro-3330-MT:~$ sudo su hduser
[sudo] password for student:
hduser@student-HP-Pro-3330-MT:/home/student$ cd ..
hduser@student-HP-Pro-3330-MT:/home$ cd ..
hduser@student-HP-Pro-3330-MT:/$ cd usr/local/hive/bin
hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ start-dfs.sh
24/04/25 10:41:25 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-
hduser-namenode-student-HP-Pro-3330-MT.out
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-
hduser-datanode-student-HP-Pro-3330-MT.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to
/usr/local/hadoop/logs/hadoop-hduser-secondarynamenode-student-HP-Pro-
3330-MT.out
24/04/25 10:41:43 WARN util.NativeCodeLoader: Unable to load native-
hadoop library for your platform... using builtin-java classes where
applicable
hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-
resourcemanager-student-HP-Pro-3330-MT.out
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-
hduser-nodemanager-student-HP-Pro-3330-MT.out
hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ jps
3792 NodeManager
4130 Jps
3122 NameNode
3492 SecondaryNameNode
3253 DataNode
3663 ResourceManager
hduser@student-HP-Pro-3330-MT:/usr/local/hive/bin$ ./hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-
2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in
[jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-
1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in
jar:file:/usr/local/hive/lib/hive-common-2.1.1.jar!/hive-
log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future
versions. Consider using a different execution engine (i.e. spark, tez)
or using Hive 1.X releases.
hive> show databases;
OK
```

```
db1
db2
db3
db74
default
flight_record
flights
hb
schedule
tanveer
tanveerkaran
test
thanos
yaduraj
Time taken: 1.402 seconds, Fetched: 14 row(s)
hive> create database stranger;
OK
Time taken: 0.832 seconds
hive> show databases;
OK
db1
db2
db3
db74
default
flight_record
flights
hb
schedule
stranger
tanveer
tanveerkaran
test
thanos
yaduraj
Time taken: 0.024 seconds, Fetched: 15 row(s)
hive> use stranger;
OK
Time taken: 0.014 seconds
hive> show tables;
OK
Time taken: 0.212 seconds
hive> create table doctor(doc_id int, doc_name string,special string);
OK
Time taken: 1.052 seconds
hive> create table patient(pat_id int, pat_name string,age int);
OK
Time taken: 0.113 seconds
hive> create table medicine(price int,qty int,name string);
OK
Time taken: 0.101 seconds
hive> insert into table doctor values(01,"Yash Varma","Neurologist"),
    > (02,"Abhishek Singh","Heart Spec."),
    > (03,"Kamlesh Kumawat","Liver Spec.");
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hduser_20240425105952_78969147-376e-4058-8cdc-94065e279f85
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714021958315_0001, Tracking URL = http://student-HP-Pro-3330-MT:8088/proxy/application_1714021958315_0001/
Kill Command = /usr/local/hadoop/bin/hadoop job _-kill job_1714021958315_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-04-25 11:00:24,104 Stage-1 map = 0%, reduce = 0%
2024-04-25 11:00:47,054 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.2 sec
MapReduce Total cumulative CPU time: 2 seconds 200 msec
Ended Job = job_1714021958315_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory
hdfs://localhost:9000/user/hive/warehouse/stranger.db/doctor/.hive-staging_hive_2024-04-25_10-59-52_691_3292171009977523598-1/-ext-10000
Loading data to table stranger.doctor
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 2.2 sec HDFS Read: 4432 HDFS Write: 155 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 200 msec
OK
Time taken: 61.135 seconds
hive> show table doctor;
MismatchedTokenException(26!=110)
at
org.antlr.runtime.BaseRecognizer.recoverFromMismatchedToken(BaseRecognizer.java:617)
at org.antlr.runtime.BaseRecognizer.match(BaseRecognizer.java:115)
at
org.apache.hadoop.hive.ql.parse.HiveParser.showStatement(HiveParser.java:20763)
at
org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:2838)
at
org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:1756)
at
org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1178)
)
at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:204)
at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)
at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:444)

```

    at
org.apache.hadoop.hive ql.Driver.compileInternal(Driver.java:1242)
    at org.apache.hadoop.hive ql.Driver.runInternal(Driver.java:1384)
    at org.apache.hadoop.hive ql.Driver.run(Driver.java:1171)
    at org.apache.hadoop.hive ql.Driver.run(Driver.java:1161)
    at
org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:232)
    at
org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:183)
    at
org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:399)
    at
org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:776)
    at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:714)
    at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:641)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java
:62)
    at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessori
mpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:239)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:153)
FAILED: ParseException line 1:11 mismatched input 'doctor' expecting
EXTENDED near 'table' in show statement
hive> desc doctor;
OK
doc_id          int
doc_name        string
special         string
Time taken: 0.129 seconds, Fetched: 3 row(s)
hive> select * from doctor;
OK
1      Yash Varma Neurologist
2      Abhishek Singh   Heart Spec.
3      Kamlesh Kumawat  Liver Spec.
Time taken: 0.281 seconds, Fetched: 3 row(s)
hive> insert into table patient values(01,"Darshan Deshmukh",21),
    > (02,"Sumit Patil",25),
    > (03,"Sanket Rajnor",32);
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = hduser_20240425110349_1ce340f2-906c-4adc-82c2-64d78ee06968
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714021958315_0002, Tracking URL = http://student-HP-
Pro-3330-MT:8088/proxy/application_1714021958315_0002/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill
job_1714021958315_0002

```

```

Hadoop job information for Stage-1: number of mappers: 1; number of
reducers: 0
2024-04-25 11:04:06,791 Stage-1 map = 0%,   reduce = 0%
2024-04-25 11:04:20,218 Stage-1 map = 100%,   reduce = 0%, Cumulative CPU
2.21 sec
MapReduce Total cumulative CPU time: 2 seconds 210 msec
Ended Job = job_1714021958315_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory
hdfs://localhost:9000/user/hive/warehouse/stranger.db/patient/.hive-
staging_hive_2024-04-25_11-03-49_762_4534713500501255151-1/-ext-10000
Loading data to table stranger.patient
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 2.21 sec   HDFS Read: 4412 HDFS
Write: 130 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 210 msec
OK
Time taken: 35.069 seconds
hive> insert into table medicine values(120,40,"Cipla X-20"),
    > (250,60,"Aerocita"),
    > (400,20,"Ostrifill");
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = hduser_20240425110642_a63aad3f-b3c8-455a-90c0-9ecc90fdbae1
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714021958315_0003, Tracking URL = http://student-HP-
Pro-3330-MT:8088/proxy/application_1714021958315_0003/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill
job_1714021958315_0003
Hadoop job information for Stage-1: number of mappers: 1; number of
reducers: 0
2024-04-25 11:06:58,934 Stage-1 map = 0%,   reduce = 0%
2024-04-25 11:07:07,250 Stage-1 map = 100%,   reduce = 0%, Cumulative CPU
2.25 sec
MapReduce Total cumulative CPU time: 2 seconds 250 msec
Ended Job = job_1714021958315_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory
hdfs://localhost:9000/user/hive/warehouse/stranger.db/medicine/.hive-
staging_hive_2024-04-25_11-06-42_039_2296046379715348999-1/-ext-10000
Loading data to table stranger.medicine
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 2.25 sec   HDFS Read: 4394 HDFS
Write: 124 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 250 msec
OK
Time taken: 28.711 seconds

```

```

hive> select * from doctor;
OK
1      Yash Varma Neurologist
2      Abhishek Singh   Heart Spec.
3      Kamlesh Kumawat Liver Spec.
Time taken: 0.129 seconds, Fetched: 3 row(s)
hive> select * from patient;
OK
1      Darshan Deshmukh 21
2      Sumit Patil 25
3      Sanket Rajnor    32
Time taken: 0.134 seconds, Fetched: 3 row(s)
hive> select * from medicine;
OK
120    40    Cipla X-20
250    60    Aerocita
400    20    Ostrifill
Time taken: 0.107 seconds, Fetched: 3 row(s)
hive> alter table patient add columns (illness string);
OK
Time taken: 0.331 seconds
hive> select * from patient
> ;
OK
1      Darshan Deshmukh 21      NULL
2      Sumit Patil 25      NULL
3      Sanket Rajnor    32      NULL
Time taken: 0.124 seconds, Fetched: 3 row(s)
hive> alter table patient where pat_id=1 set illness="Fever";
NoViableAltException(305@[])
    at
org.apache.hadoop.hive.ql.parse.HiveParser.alterTableStatementSuffix(Hive
Parser.java:8010)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.alterStatement(HiveParser.java
:7446)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:2
808)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:
1756)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1178
)
    at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:204)
    at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)
    at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:444)
    at
org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1242)
    at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1384)
    at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1171)

```

```

        at org.apache.hadoop.hive ql.Driver.run(Driver.java:1161)
        at
org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:232)
        at
org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:183)
        at
org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:399)
        at
org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:776)
        at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:714)
        at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:641)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java
:62)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessori
mpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:239)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:153)
FAILED: ParseException line 1:20 cannot recognize input near 'where'
'pat_id' '=' in alter table statement
hive> alter table patient where pat_id = 1 set illness="Fever";
NoViableAltException(305@[])
        at
org.apache.hadoop.hive ql.parse.HiveParser.alterTableStatementSuffix(Hive
Parser.java:8010)
        at
org.apache.hadoop.hive ql.parse.HiveParser.alterStatement(HiveParser.java
:7446)
        at
org.apache.hadoop.hive ql.parse.HiveParser.ddlStatement(HiveParser.java:2
808)
        at
org.apache.hadoop.hive ql.parse.HiveParser.execStatement(HiveParser.java:
1756)
        at
org.apache.hadoop.hive ql.parse.HiveParser.statement(HiveParser.java:1178
)
        at
org.apache.hadoop.hive ql.parse.ParseDriver.parse(ParseDriver.java:204)
        at
org.apache.hadoop.hive ql.parse.ParseDriver.parse(ParseDriver.java:166)
        at org.apache.hadoop.hive ql.Driver.compile(Driver.java:444)
        at
org.apache.hadoop.hive ql.Driver.compileInternal(Driver.java:1242)
        at org.apache.hadoop.hive ql.Driver.runInternal(Driver.java:1384)
        at org.apache.hadoop.hive ql.Driver.run(Driver.java:1171)
        at org.apache.hadoop.hive ql.Driver.run(Driver.java:1161)
        at
org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:232)
        at
org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:183)

```

```

        at
org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:399)
        at
org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:776)
        at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:714)
        at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:641)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java
:62)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessori
mpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:239)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:153)
FAILED: ParseException line 1:20 cannot recognize input near 'where'
'pat_id' '=' in alter table statement
hive> update patient set illness="Fever" where pat_id=1;;
FAILED: SemanticException [Error 10294]: Attempt to do update or delete
using transaction manager that does not support these operations.
hive> update patient set illness="Fever" where pat_id=1;
FAILED: SemanticException [Error 10294]: Attempt to do update or delete
using transaction manager that does not support these operations.
hive> select d.doc_id, doc_name, p.pat_name, p.age from doctor d join
patient p on (d.doc_id=p.pat_id);
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = hduser_20240425111815_02da7588-ef51-4a4d-af6d-caeb1b74527d
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-
2.4.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in
[jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-
1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFactory]
2024-04-25 11:18:28      Starting to launch local task to process map join;
        maximum memory = 477626368
2024-04-25 11:18:30      Dump the side-table for tag: 1 with group count: 3
into file: file:/tmp/hduser/a33832f1-8f33-4330-9db4-
7932f38b748e/hive_2024-04-25_11-18-15_348_3796776716572301718-1/-local-
10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable
2024-04-25 11:18:30      Uploaded 1 File to: file:/tmp/hduser/a33832f1-
8f33-4330-9db4-7932f38b748e/hive_2024-04-25_11-18-
15_348_3796776716572301718-1/-local-10004/HashTable-Stage-3/MapJoin-
mapfile01--.hashtable (363 bytes)
2024-04-25 11:18:30      End of local task; Time Taken: 2.12 sec.
Execution completed successfully
MapredLocal task succeeded

```



```

Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714021958315_0004, Tracking URL = http://student-HP-
Pro-3330-MT:8088/proxy/application_1714021958315_0004/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill
job_1714021958315_0004
Hadoop job information for Stage-3: number of mappers: 1; number of
reducers: 0
2024-04-25 11:18:45,195 Stage-3 map = 0%, reduce = 0%
2024-04-25 11:18:55,628 Stage-3 map = 100%, reduce = 0%, Cumulative CPU
2.31 sec
MapReduce Total cumulative CPU time: 2 seconds 310 msec
Ended Job = job_1714021958315_0004
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 2.31 sec HDFS Read: 6861 HDFS
Write: 223 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 310 msec
OK
1 Yash Varma Darshan Deshmukh 21
2 Abhishek Singh Sumit Patil 25
3 Kamlesh Kumawat Sanket Rajnor 32
Time taken: 43.5 seconds, Fetched: 3 row(s)
hive> create index cidx on table medicine(name) as
'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' with
deferred rebuild;
MismatchedTokenException(26!=86)
    at
org.antlr.runtime.BaseRecognizer.recoverFromMismatchedToken(BaseRecognize
r.java:617)
    at org.antlr.runtime.BaseRecognizer.match(BaseRecognizer.java:115)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.autoRebuild(HiveParser.java:64
64)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.createIndexStatement(HiveParse
r.java:6007)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:2
928)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:
1756)
    at
org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1178
)
    at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:204)
    at
org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)
    at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:444)
    at
org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1242)
    at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1384)
    at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1171)

```

```
    at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1161)
    at
org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:232)
    at
org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:183)
    at
org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:399)
    at
org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:776)
    at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:714)
    at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:641)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java
:62)
    at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessori
mpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:239)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:153)
```

FAILED: ParseException line 1:112 mismatched input 'deferred' expecting
KW_DEFERRED near 'with' in auto rebuild index

```
hive> create index cidx on table medicine(name) as
'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' with
deferred rebuild;
```

OK

Time taken: 0.622 seconds

```
hive> select avg(price) from medicine;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.

Query ID = hduser_20240425112310_24d325d9-acf5-4271-b0c9-dd3dbb68043a

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
    set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
    set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
    set mapreduce.job.reduces=<number>
```

Starting Job = job_1714021958315_0005, Tracking URL = http://student-HP-
Pro-3330-MT:8088/proxy/application_1714021958315_0005/

Kill Command = /usr/local/hadoop/bin/hadoop job -kill

job_1714021958315_0005

Hadoop job information for Stage-1: number of mappers: 1; number of
reducers: 1

2024-04-25 11:23:29,261 Stage-1 map = 0%, reduce = 0%

2024-04-25 11:23:42,368 Stage-1 map = 100%, reduce = 0%, Cumulative CPU
2.03 sec

2024-04-25 11:24:03,878 Stage-1 map = 100%, reduce = 100%, Cumulative
CPU 4.84 sec

MapReduce Total cumulative CPU time: 4 seconds 840 msec

```
Ended Job = job_1714021958315_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.84 sec HDFS Read:
8596 HDFS Write: 117 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 840 msec
OK
256.6666666666667
Time taken: 56.423 seconds, Fetched: 1 row(s)
hive> show index on medicine
> ;
OK
cidx          medicine          name
stranger__medicine_cidx__ compact
Time taken: 0.166 seconds, Fetched: 1 row(s)
hive> exit;
```